

A STUDY OF RANDOM MATRICES

Sample Covariance Matrices and Spectral Clustering

AN INDEPENDENT STUDY REPORT

SUBMITTED IN [MONTH], 2026
TO THE DEPARTMENT OF MATHEMATICS
OF NEWCOMB-TULANE COLLEGE
TULANE UNIVERSITY

BY

Arzaan Singh

UNDER THE SUPERVISION OF

Gustavo Didier, Ph.D.
Department of Mathematics

Abstract

[TODO] Write the abstract once the body is finalized. Suggested skeleton: (1) Open with a one-sentence framing of random matrix theory and why sample covariance matrices matter in high-dimensional statistics. (2) State the scope: we focus on Gaussian sample covariance matrices in Johnstone’s spiked model. (3) Outline the journey: foundational probability and linear algebra, sample covariance matrices and the Wishart distribution, low-dimensional fluctuation behavior, the high-dimensional Marchenko-Pastur and Tracy-Widom laws, the spiked covariance model and the BBP phase transition, eigenvector alignment beyond the threshold, and a sketch of the connection to spectral clustering. (4) Note the role of Monte Carlo simulations in illustrating each phase of the argument. (5) Close with the takeaway: the contrast between strong and weak spikes produces a clean phase transition in both eigenvalue location and eigenvector informativeness, with implications for spectral clustering. Avoid em dashes; use clean academic prose.

Acknowledgements

[TODO] *Acknowledgements paragraph. Suggested mentions: Professor Gustavo Didier (advisor) for the weekly meetings, the reading recommendations (Yao, Zheng, and Bai 2015; Féral and Pécché 2009), and the patient guidance through the spiked covariance literature. The second reader (TBD) for serving on the committee. Optional: family, friends, and the Tulane Department of Mathematics.*

Contents

Abstract	ii
Acknowledgements	iii
1 Introduction	1
2 Probability Foundations	4
2.1 Random Variables and Distributions	4
2.2 Moments and Moment Generating Functions	4
2.3 Characteristic Functions	5
2.4 Modes of Convergence	5
2.5 Limit Theorems: LLN and CLT	6
2.6 Random Vectors and the Multivariate Normal Distribution	6
3 Sample Covariance Matrices	8
3.1 Population and Sample Covariance, Spectral Theorem, Positive Semidefiniteness	8
3.2 Low Dimensions: Wishart, Eigenvalue Repulsion, and Fluctuations	8
3.3 High Dimensions: Marchenko-Pastur and the Tracy-Widom Edge	9
4 Spiked Covariance Matrices and the BBP Transition	10
4.1 Johnstone’s Spiked Model and the Spike-Location Map	10
4.2 The BBP Phase Transition for Eigenvalues	10
4.3 Fluctuations of Spiked Eigenvalues	11
4.4 Eigenvector Alignment Beyond the Threshold	11
4.5 Sketch of the Connection to Spectral Clustering	11
5 Conclusion and Outlook	12
A Auxiliary Results	13
B Simulation Methodology	14
C Selected Code Listings	15

1 Introduction

Why random matrices?

Modern data sets are large and high-dimensional. A single-cell genomics experiment may measure tens of thousands of gene expression levels in a few hundred cells. A finance application may track hundreds of asset returns over a few years of daily data. A computer-vision system may compare millions of image patches in a feature space of several hundred dimensions. In each case the number of variables p is comparable to, or larger than, the number of observations n . The most basic question one can ask of such data is how its variance is distributed across orthogonal directions, and which directions explain the most variation. Answering this question is the job of *principal components analysis*, and the algorithm reduces to computing the eigenvalues and eigenvectors of the *sample covariance matrix*

$$\hat{\Sigma}_n = \frac{1}{n} \sum_{i=1}^n \mathbf{X}_i \mathbf{X}_i^\top, \quad (1)$$

where the $\mathbf{X}_i \in \mathbb{R}^p$ are the n observed data points, assumed centered.

Classical multivariate statistics tells us what to expect when p is small and n is large. In this low-dimensional regime, $\hat{\Sigma}_n$ is an accurate estimate of the population covariance Σ , its eigenvalues converge to the population eigenvalues, and its eigenvectors converge to the population eigenvectors. Methods that depend on this convergence, from PCA to factor analysis to the *spectral clustering* algorithms that we will reach at the end of this report, then stand on firm ground.

The classical picture breaks down when p and n are comparable. The sample covariance matrix is no longer a reliable estimate of Σ : its eigenvalues spread into a non-trivial band of random values, and its eigenvectors may carry little or no information about the population directions. Figure 1 makes the contrast concrete. The same Gaussian sampling model produces eigenvalues that concentrate at the population value 1 when $p = 10$ and $n = 1000$, and eigenvalues that spread across a wide interval, with a well-defined limiting density, when $p = 500$ and $n = 1000$.

The mathematical framework that explains and predicts this phenomenon is *random matrix theory*, originally developed for problems in nuclear physics in the 1950s and now a central tool in high-dimensional statistics, signal processing, and machine learning (Yao et al., 2015; Vershynin, 2018; Couillet and Liao, 2022).

What this report does, and how

This report develops, from first principles, the spectral analysis of sample covariance matrices in both regimes, and uses what we learn to understand when methods like spectral clustering can recover hidden structure in high-dimensional data. The development is built bottom-up: every claim is either proved in the text, sketched with a pointer to a full proof, or cited with the cleanest reference we could find. Every theoretical claim is then illustrated with a Monte Carlo simulation in Python; the notebooks live in the `notebooks/` directory of




Figure 1: Eigenvalues of the sample covariance matrix $\hat{\Sigma}_n$ when $\Sigma = I_p$ and the data are Gaussian, for two aspect ratios. *Left:* when $p = 10$ and $n = 1000$ (so $p/n = 0.01$), the ten sample eigenvalues concentrate tightly around the population value $\lambda = 1$, confirming the classical low-dimensional intuition. *Right:* when $p = 500$ and $n = 1000$ (so $p/n = 0.5$), the five hundred sample eigenvalues spread across the interval $[(1 - \sqrt{y})^2, (1 + \sqrt{y})^2] \approx [0.086, 2.914]$. The empirical density agrees with the Marchenko-Pastur prediction (solid curve). The phenomenon that motivates this report is already visible: in high dimensions, sample eigenvalues are biased, and the spread is a real feature of the geometry, not measurement noise.

the accompanying repository, and the figures they produce are reproduced alongside the corresponding statements.

The starting point is the random variable, the basic building block of probability theory. From there we build to moments and moment generating functions, characteristic functions, modes of convergence, the central limit theorem, random vectors with multivariate normal distributions, and the spectral theorem for symmetric matrices. Each of these objects feeds directly into the analysis of $\hat{\Sigma}_n$ that follows.

The story arc

The mathematical journey unfolds in three movements.

1. *The low-dimensional regime* (p fixed, $n \rightarrow \infty$). The law of large numbers gives entrywise convergence of $\hat{\Sigma}_n$ to Σ , and continuity of the characteristic polynomial then forces the sample eigenvalues to converge to the population eigenvalues. The size and shape of the fluctuations around these limits depend on whether the population eigenvalues are simple or repeated. For $\Sigma = I_p$ the matrix $\sqrt{n}(\hat{\Sigma}_n - I_p)$ converges in distribution to a *Gaussian Orthogonal Ensemble* matrix, whose top eigenvalue has a non-Gaussian limit shaped by *eigenvalue repulsion*. For Σ with distinct eigenvalues, the delta method instead produces ordinary Gaussian fluctuations of each sample eigenvalue around its population counterpart.
2. *The high-dimensional regime* ($p, n \rightarrow \infty$ with $p/n \rightarrow y \in (0, \infty)$). Even for the null covariance $\Sigma = I_p$, the sample eigenvalues no longer concentrate at 1; they fill out the *Marchenko-Pastur law* on the interval $[(1 - \sqrt{y})^2, (1 + \sqrt{y})^2]$ (Marchenko and Pastur,

1967), exactly as in the right panel of Figure 1. The largest eigenvalue converges almost surely to the upper edge $(1 + \sqrt{y})^2$, but its fluctuations around that edge are not Gaussian. They live on the scale $p^{-2/3}$ and follow the *Tracy-Widom law* (Tracy and Widom, 1996; Johnstone, 2001), a distinct universality class first discovered in random matrix theory and now found in problems ranging from longest increasing subsequences to growth models in mathematical physics.

3. *Structure and a phase transition.* When the population covariance carries low-rank structure, a small number of distinguished eigenvalues riding above a flat baseline, a striking new phenomenon appears. If a signal eigenvalue exceeds a threshold determined by the aspect ratio y , an isolated outlier eigenvalue appears in the sample spectrum and the corresponding sample eigenvector aligns nontrivially with the population spike direction. If the signal is below the threshold, the spike is absorbed into the Marchenko-Pastur bulk and its eigenvector becomes asymptotically orthogonal to the population direction. This phase transition is the bridge to spectral clustering: the same picture explains how clustering algorithms recover hidden class structure from data, but only when the underlying signal-to-noise ratio crosses an analogous threshold.

Where this is applicable

The techniques developed here are not abstract curiosities. The Marchenko-Pastur law sets a baseline for detecting structure in any high-dimensional covariance estimate, with applications in array-signal processing, wireless communications, and quantitative finance (Couillet and Liao, 2022). The Tracy-Widom edge governs the null distribution of the largest principal component, which is the basis for principled testing of how many genuine signal directions are present in a data set (Johnstone, 2001). The spike phase transition determines when a low-rank signal in noise is recoverable by PCA, with direct implications for denoising and dimensionality reduction. Spectral clustering, finally, is one of the most widely used unsupervised methods for problems where the natural geometry is not Euclidean, including image segmentation, community detection in networks, and clustering of single-cell data (von Luxburg, 2007; Ng et al., 2002). The phase transition we develop in the spiked model is the cleanest mathematical lens we have for understanding when these methods succeed and when they fail.

Roadmap

Section 2 reviews probability foundations: random variables, moments and moment generating functions, characteristic functions, modes of convergence, the law of large numbers and central limit theorem, and random vectors. Section 3 develops the sample covariance matrix in both regimes, from the spectral theorem and the Wishart distribution in low dimensions through the Marchenko-Pastur and Tracy-Widom laws in high dimensions. Section 4 introduces the spiked covariance model, proves the eigenvalue phase transition, derives the eigenvector overlap formula, and sketches the connection to spectral clustering. Section 5 concludes with a summary and open questions, and three appendices collect auxiliary probabilistic results, simulation methodology, and selected Python code listings.

2 Probability Foundations

[TODO] *Opening paragraph: motivation. The high-dimensional phenomena we will study in later sections (Marchenko-Pastur, Tracy-Widom, the spike phase transition) all rest on a small set of foundational ideas from classical probability. This section develops them in depth, starting from the definition of a random variable. The depth here is deliberate: most of the later analysis is, at its heart, a careful application of these basic tools (differentiating under the integral, characteristic function expansions, the multivariate central limit theorem, the spectral theorem). A reader who is fluent with this material can skim; everyone else benefits from the deliberate pace.*

2.1 Random Variables and Distributions

[TODO] *Build absolute foundation (from v5 + Joseph Section 2.1):*

- *Definition: probability space $(\Omega, \mathcal{F}, \mathbb{P})$ at a light touch.*
- *Definition: random variable $X : \Omega \rightarrow \mathbb{R}$.*
- *Definition: CDF $F_X(x) = \mathbb{P}(X \leq x)$ and its three characterizing properties (limits, right-continuity, monotonicity). State the converse (any function satisfying these is a CDF), cite (Casella and Berger, 2002, §1.2).*
- *Definition: discrete vs. absolutely continuous random variables; PMF and PDF.*
- *Examples: Bernoulli, binomial, Poisson, normal, gamma, Cauchy. Densities listed.*
- *Definition: independence of random variables.*

Cite Casella-Berger and Joseph throughout.

2.2 Moments and Moment Generating Functions

[TODO] *Full depth from v5. Sub-structure:*

- *Definition: j th moment, j th central moment, variance.*
- *Definition: MGF $M_X(t) = \mathbb{E}[e^{tX}]$ with existence requirement on an open neighborhood of zero.*
- *Heuristic: MGF as the two-sided Laplace transform of the density (use this to motivate uniqueness).*
- *Theorem: if M_X exists in a neighborhood of 0, then $M_X^{(n)}(0) = \mathbb{E}[X^n]$. Full proof by differentiation under the integral (continuous case) and term-by-term differentiation (discrete case).*

- *Theorem: MGF uniqueness. If $M_X(t) = M_Y(t)$ on a neighborhood of 0, then $X \stackrel{d}{=} Y$. State precisely; sketch the easy direction; cite (Casella and Berger, 2002, Thm. 2.3.11) for the converse.*
- *Worked examples: MGFs of binomial, Poisson, normal, gamma. Derive each. Tabulate the resulting means and variances.*
- *Cautionary example: Cauchy has no MGF (the integral diverges for all $t \neq 0$). Use this to motivate characteristic functions in §2.3.*

2.3 Characteristic Functions

[TODO] Full depth from v5. Sub-structure:

- *Definition: $\varphi_X(t) = \mathbb{E}[e^{itX}]$. Always exists since $|e^{itX}| = 1$; deduce $|\varphi_X(t)| \leq 1$ and $\varphi_X(0) = 1$.*
- *Identification as a Fourier transform of the density (when one exists), and as a Fourier-Stieltjes transform of the distribution in general.*
- *Relation to the MGF when both exist: $\varphi_X(t) = M_X(it)$.*
- *Worked examples: derive φ_X for the standard normal (completing the square gives $e^{-t^2/2}$) and the Cauchy (gives $e^{-|t|}$). The Cauchy case is the punchline: characteristic functions handle heavy tails that defeat MGFs.*
- *Properties: independence implies $\varphi_{X+Y}(t) = \varphi_X(t)\varphi_Y(t)$; linear transformations $\varphi_{aX+b}(t) = e^{itb}\varphi_X(at)$.*
- *Statement of the uniqueness theorem for CFs (the Fourier inversion content).*
- *Statement of Lévy's continuity theorem: pointwise convergence of CFs to a function continuous at 0 is equivalent to convergence in distribution. Defer the full proof to Appendix A; cite (Casella and Berger, 2002, Thm. 5.5.15).*

Suggested figure: side-by-side plot of $\varphi_{N(0,1)}(t) = e^{-t^2/2}$ and $\varphi_{\text{Cauchy}}(t) = e^{-|t|}$, to make tangible the contrast between light-tailed and heavy-tailed behavior at the CF level.

2.4 Modes of Convergence

[TODO] Content (from v5 + Joseph Section 2.2):

- *Definitions: $X_n \xrightarrow{\text{a.s.}} X$ (almost-sure convergence), $X_n \xrightarrow{\mathbb{P}} X$ (convergence in probability), $X_n \xrightarrow{d} X$ (convergence in distribution).*
- *Theorem: implication chain a.s. $\Rightarrow$ probability $\Rightarrow$ distribution. Prove the first implication via the Borel-Cantelli lemma; sketch the second.*

- Counterexamples to the converses:
 - $X_n \in \{0, 1\}$ with $\mathbb{P}(X_n = 1) = 1/n$ converges in probability to 0 but not almost surely (the second Borel-Cantelli lemma forces $X_n = 1$ infinitely often).
 - For convergence in distribution but not in probability, use $X \sim \mathcal{N}(0, 1)$ and $X_n = -X$ (the distribution is symmetric).
- Continuous mapping theorem: if $X_n \xrightarrow{a.s.} X$ and g is continuous, then $g(X_n) \xrightarrow{a.s.} g(X)$; same for convergence in probability. Defer Slutsky's theorem to Appendix A.

2.5 Limit Theorems: LLN and CLT

[TODO] Content (from v5 + Joseph Section 2.2):

- Inequalities first: Markov, then Chebyshev (the latter as a Markov application to $(X - \mu)^2$). Full proofs.
- Chebyshev's weak law of large numbers: iid with finite variance $\Rightarrow \bar{X}_n \xrightarrow{\mathbb{P}} \mu$. Full proof via Chebyshev.
- Statement of Kolmogorov's strong law of large numbers; cite (Billingsley, 1995) for the proof.
- Central limit theorem: for iid with finite variance, $\sqrt{n}(\bar{X}_n - \mu)/\sigma \xrightarrow{d} \mathcal{N}(0, 1)$. Prove via characteristic function expansion using the Lévy continuity theorem from §2.3. Carry out the $(1 - t^2/(2n) + o(1/n))^n \rightarrow e^{-t^2/2}$ step explicitly.

Suggested figure (notebook 01): the WLLN convergence for normal and exponential samples, plotting $\bar{X}_n$ versus n . Reproduce Joseph's Figure 1 in our color palette. Also a second figure showing the CLT convergence: standardized sample means at $n = 10, 100, 1000$ versus the standard normal density.

2.6 Random Vectors and the Multivariate Normal Distribution

[TODO] Content (from v5 + Joseph Section 2.4):

- Definition: random vector $\mathbf{X} = (X_1, \dots, X_p)^\top$, joint CDF, joint PMF / PDF.
- Mean vector $\boldsymbol{\mu} = \mathbb{E}[\mathbf{X}]$ and covariance matrix $\Sigma = \mathbb{E}[(\mathbf{X} - \boldsymbol{\mu})(\mathbf{X} - \boldsymbol{\mu})^\top]$.
- Theorem: Σ is positive semidefinite. Full proof via $v^\top \Sigma v = \text{Var}(v^\top \mathbf{X}) \geq 0$.
- When is Σ positive definite vs. semidefinite? Connect to the dimension of the support.
- Definition: multivariate normal distribution via density. State the equivalent characterization via characteristic function $\varphi_{\mathbf{X}}(\mathbf{t}) = \exp(i\mathbf{t}^\top \boldsymbol{\mu} - \frac{1}{2}\mathbf{t}^\top \Sigma \mathbf{t})$; this is the cleanest definition because it handles singular Σ without modification.

- *Properties: linear transformations of multivariate normals are multivariate normal; uncorrelated jointly Gaussian variables are independent (special fact for the Gaussian).*
- *Statement of the multivariate central limit theorem; the proof is the same Cramér-Wold device applied to the univariate CLT, defer to Appendix [A](#).*

3 Sample Covariance Matrices

[TODO] Opening paragraph: the sample covariance matrix $\hat{\Sigma}_n$ is the central object of multivariate statistics. In low dimensions (p fixed, $n \rightarrow \infty$) it converges almost surely to the population covariance Σ and its eigenvalues converge by continuity. In high dimensions ($p, n \rightarrow \infty$ with $p/n \rightarrow y > 0$) this convergence fails dramatically: the empirical spectrum spreads out into a non-trivial bulk density. This section develops both regimes, building from the foundations laid in Section 2.

3.1 Population and Sample Covariance, Spectral Theorem, Positive Semidefiniteness

[TODO] Content to fill in (from v5 + Joseph Section 3 intro): - Definition: sample covariance matrix $\hat{\Sigma}_n = \frac{1}{n} \sum_{i=1}^n \mathbf{X}_i \mathbf{X}_i^\top$ (mean-zero case) and the unbiased $\frac{1}{n-1}$ variant. - Matrix form: $\hat{\Sigma}_n = \frac{1}{n} \mathbf{X} \mathbf{X}^\top$ for $\mathbf{X} \in \mathbb{R}^{p \times n}$. - Spectral theorem for symmetric matrices: $A = Q \Lambda Q^\top$. - PSD characterization: A is PSD iff all eigenvalues are nonneg (proof). - Use the diagonalization to compute quadratic forms; foreshadow why eigenvalues of $\hat{\Sigma}_n$ are the meaningful quantities.

3.2 Low Dimensions: Wishart, Eigenvalue Repulsion, and Fluctuations

[TODO] Content to fill in (from v5 + Joseph Section 3.1):

(a) Almost sure convergence of the sample covariance. - Theorem: if $\mathbf{X}_i$ iid with $\mathbb{E} \mathbf{X}_1 = 0$ and $\text{Cov}(\mathbf{X}_1) = \Sigma$, then $\hat{\Sigma}_n \xrightarrow{\text{a.s.}} \Sigma$ entrywise by SLLN. - Corollary: $\lambda_i(\hat{\Sigma}_n) \xrightarrow{\text{a.s.}} \lambda_i(\Sigma)$ by continuity of eigenvalues (proof sketch via continuity of the characteristic polynomial and continuity of its roots). Cite Genzer (2025, Prop. 3.1, 3.3).

(b) Wishart distribution. - Definition: $\mathbf{S} = \sum_{i=1}^n \mathbf{X}_i \mathbf{X}_i^\top \sim W_p(n, \Sigma)$ for iid Gaussian $\mathbf{X}_i \sim \mathcal{N}(0, \Sigma)$. - Wishart density formula. - Theorem: for $p = 1$, $W_1(n, 1)$ collapses to χ_n^2 (proof by plugging in).

(c) 2D Wishart and eigenvalue repulsion. - Joint density of ordered eigenvalues of $W_2(n, I_2)$ contains the factor $\lambda_2 - \lambda_1$, producing eigenvalue repulsion. - Simulation: scatter of (λ_1, λ_2) and histogram of $\lambda_2 - \lambda_1$ (notebook 02). Include figure: 2D Wishart eigenvalue scatter.

(d) Joint CLT for entries of $\hat{\Sigma}_n$ when $\Sigma = I_2$. - Statement and proof via multivariate CLT (variances 2, 2, 1 for diagonal and off-diagonal entries). - Consequence: matrix-valued CLT $\sqrt{n}(\hat{\Sigma}_n - I_2) \xrightarrow{d} G$ where G has GOE structure.

(e) Top eigenvalue fluctuation: repeated vs. distinct cases. - Repeated ($\Sigma = I_2$): $\sqrt{n}(\lambda_2(\hat{\Sigma}_n) - 1) \xrightarrow{d} \lambda_2(G)$, not Gaussian. - Distinct ($\Sigma = \text{diag}(1, 2)$): delta-method gives Gaussian fluctuations with variances 2 and 8. - Simulation: Q-Q plots for both cases (notebook 03). Include figures: Q-Q repeated, Q-Q distinct.

(f) *Eigenvector convergence.* - Distinct case: $\widehat{\mathbf{u}}_\ell \xrightarrow{\text{a.s.}} \mathbf{e}_\ell$ with sign convention. - Repeated case: eigenvectors fail to converge; explicit counterexample with rotating O_n . This previews why spectral separation is essential for eigenvector recovery, a theme that returns in Section 4.

3.3 High Dimensions: Marchenko-Pastur and the Tracy-Widom Edge

[TODO] Content to fill in (from v5 high-dim section + Joseph Section 3.2):

(a) *The high-dimensional regime.* - Setup: $p, n \rightarrow \infty$ with $p/n \rightarrow y \in (0, \infty)$. - Empirical spectral distribution $F_{\widehat{\Sigma}_n}(\lambda) = \frac{1}{p} \#\{i : \lambda_i(\widehat{\Sigma}_n) \leq \lambda\}$.

(b) *Marchenko-Pastur law.* - Theorem (Marchenko-Pastur 1967): for null $\Sigma = I_p$ and $0 < y < 1$, $F_{\widehat{\Sigma}_n} \xrightarrow{d} \text{MP}_y$ with support $[(1 - \sqrt{y})^2, (1 + \sqrt{y})^2]$ and density $f_{\text{MP}}(\lambda) = \frac{1}{2\pi y \lambda} \sqrt{(\lambda_+ - \lambda)(\lambda - \lambda_-)}$. - For $y \geq 1$, an atom at zero of mass $1 - 1/y$. - Convergence of $\lambda_{\max}(\widehat{\Sigma}_n) \xrightarrow{\text{a.s.}} (1 + \sqrt{y})^2$. - Simulation: histogram against MP density at $y = 1/2$ (notebook 04). Include figure: MP density vs. simulation.

(c) *Tracy-Widom edge heuristic.* - Goal: explain why edge fluctuations live on scale $p^{-2/3}$, not $n^{-1/2}$. - Heuristic: square-root density near the edge, $f_{\text{MP}}(\lambda) \sim C(\lambda_+ - \lambda)^{1/2}$ as $\lambda \uparrow \lambda_+$, combined with the requirement that the mass to the right of $\lambda_{\max}$ be of order $1/p$, yields $\lambda_+ - \lambda_{\max} \sim p^{-2/3}$. - Statement (informal): $p^{2/3}(\lambda_{\max} - \lambda_+) \xrightarrow{d} \text{TW}_1$ after model-dependent constants. - Remark on large-deviation heuristic confirming the same scaling from left and right tails. Cite Tracy and Widom (1996); Johnstone (2001).

4 Spiked Covariance Matrices and the BBP Transition

[TODO] Opening paragraph: the Marchenko-Pastur law of Section 3.3 described the null case $\Sigma = I_p$, where all population eigenvalues are equal. In applications the population covariance carries structure: a small number of distinguished directions of high variance riding above a flat noise baseline. The spiked covariance model of [Johnstone \(2001\)](#) formalizes this,

$$\Sigma = \text{diag}(\alpha_1, \dots, \alpha_m, 1, \dots, 1),$$

with m fixed as $p \rightarrow \infty$. The central question is when these population spikes are detectable from the sample spectrum, and when the sample eigenvectors carry meaningful information about the population spike directions. The answer is a clean phase transition.

4.1 Johnstone's Spiked Model and the Spike-Location Map

[TODO] Content to fill in (from v5 + Yao-Zheng-Bai Ch. 11 + Gustavo notes): - Definition: Johnstone's spiked model $\Sigma = \text{diag}(\alpha_1, \dots, \alpha_m, 1, \dots, 1)$ with $H = \delta_1$ for the bulk limiting distribution. - Definition: spike-location map $\Psi(\alpha) = \alpha + \frac{y\alpha}{\alpha-1}$. - Compute $\Psi'(\alpha) = 1 - \frac{y}{(\alpha-1)^2}$. - Lemma: $\Psi'(\alpha) > 0 \iff |\alpha - 1| > \sqrt{y}$; for large spikes $\alpha > 1$ this becomes $\alpha > 1 + \sqrt{y}$. - Verify boundary behavior: $\Psi(1 \pm \sqrt{y}) = (1 \pm \sqrt{y})^2$, so the spike map agrees with the MP edges at the threshold.

4.2 The BBP Phase Transition for Eigenvalues

[TODO] Content to fill in: - Theorem (Yao-Zheng-Bai Corollary 11.4 specialized to Johnstone): For a large spike $\alpha_j > 1 + \sqrt{y}$,

$$\lambda_i(\widehat{\Sigma}_n) \xrightarrow{\text{a.s.}} \Psi(\alpha_j), \quad i \in J_j.$$

For $1 < \alpha_j \leq 1 + \sqrt{y}$,

$$\lambda_i(\widehat{\Sigma}_n) \xrightarrow{\text{a.s.}} (1 + \sqrt{y})^2.$$

Brief sketch (cite [Baik and Silverstein \(2006\)](#) for the proof). - Worked example: $y = 1/2$, so $1 + \sqrt{y} \approx 1.707$ and edge $(1 + \sqrt{y})^2 \approx 2.914$. For $\alpha = 2.5$, $\Psi(2.5) \approx 3.333$. For $\alpha = 1.4$, $\alpha < 1.707$ so the spike is absorbed into the edge. - Simulations (existing figures from v5): Figure: smoothed $m=1$ simulation, strong vs. weak spike ([smoothed_m1_strong_vs_weak.png](#)). Figure: raw top eigenvalue distributions ([m1_strong_spike.png](#), [m1_weak_spike.png](#)). Figure (optional): $m = 2$ cases ([m2_two_strong_spikes.png](#), [m2_mixed_spikes.png](#)). Cite [Baik et al. \(2005\)](#); [Yao et al. \(2015\)](#).

4.3 Fluctuations of Spiked Eigenvalues

[TODO] Content to fill in (from Gustavo's promised sketch + v5): - Statement (fundamental spike, qualitative): for $\alpha > 1 + \sqrt{y}$ in the Gaussian Johnstone model,

$$\sqrt{n}(\lambda_{\max}(\hat{\Sigma}_n) - \Psi(\alpha)) \xrightarrow{d} \mathcal{N}(0, \sigma_{\alpha,y}^2).$$

Cite Yao et al. (2015, Thm. 11.11) (fundamental spikes) and Féral and Péché (2009, Thm. 1.6) (cleanest diagonal case). - Heuristic: a separated eigenvalue is a smooth functional of $\hat{\Sigma}_n$, so a delta-method argument gives $\sqrt{n}$ Gaussian fluctuations (compare with Section 3.2(e)). - Non-fundamental spike: top eigenvalue attaches to the edge, $\lambda_{\max} - \lambda_+ = O_p(p^{-2/3})$, and the fluctuation is expected to be Tracy-Widom type (cite Johnstone (2001)). - Honest note: the critical case $\alpha = 1 + \sqrt{y}$ is not resolved in our current treatment; cite open literature. - Simulation: histogram and Gaussian Q-Q for strong-spike fluctuations (notebook 06).

4.4 Eigenvector Alignment Beyond the Threshold

[TODO] Content to fill in (from v5 + Yao-Zheng-Bai Thm 11.5/Cor 11.6): - Set-up: in the one-spike model, the population spike direction is $\mathbf{e}_1$. Let $\hat{\mathbf{u}}_1$ be the top sample eigenvector. The squared overlap $|\langle \hat{\mathbf{u}}_1, \mathbf{e}_1 \rangle|^2 = \cos^2 \theta$ removes the sign ambiguity. - Qualitative theorem (Yao-Zheng-Bai Thm. 11.5, Cor. 11.6): strong spike $\Rightarrow$ overlap converges to a positive constant; weak spike $\Rightarrow$ overlap converges to zero. - The one-spike Johnstone overlap formula:

$$\rho(\alpha, y) = \frac{1 - \frac{y}{(\alpha - 1)^2}}{1 + \frac{y}{\alpha - 1}}.$$

Positive precisely when $\alpha > 1 + \sqrt{y}$. - Simulation (figure already in v5/figures/): **eigenvector_overlap_dist** shows strong-spike overlaps clustered around 0.55–0.65 and weak-spike overlaps near 0. Include and discuss.

4.5 Sketch of the Connection to Spectral Clustering

[TODO] Content to fill in (from Gustavo's email + v5 + von Luxburg 2007): - Setup: data points $\mathbf{x}_1, \dots, \mathbf{x}_n \in \mathbb{R}^p$, kernel similarity $K_{ij} = f(\|\mathbf{x}_i - \mathbf{x}_j\|^2/p)$. - Definition: degree matrix $D = \text{diag}(K\mathbf{1})$ and normalized Laplacian-style kernel $L = nD^{-1/2}KD^{-1/2}$. - Heuristic decomposition: $L = \text{noise bulk} + \text{low-rank class signal}$, in direct analogy with $\hat{\Sigma}_n = \text{noise} + \text{spike}$. - Conclusion: class structure is recoverable from L precisely when the signal-to-noise ratio crosses an analogue of the BBP threshold; the informative eigenvectors of L then provide the clustering. Cite von Luxburg (2007) for the algorithmic side and Couillet and Benaych-Georges (2016) for the high-dimensional analysis. - Toy simulation: notebook 08, a two-class Gaussian mixture and the top non-trivial eigenvector of L .

5 Conclusion and Outlook

[TODO] Content to fill in: (a) Summary table: the strong vs. weak spike contrast, three rows: eigenvalue limit, eigenvector overlap, fluctuation scale.

— Spike strength — Eigenvalue limit — Overlap with $\mathbf{e}_1$ — Fluctuation — ————
 — Strong $\alpha > 1 + \sqrt{y}$ — $\Psi(\alpha)$
 — $\rho(\alpha, y) > 0$ — Gaussian, $\sqrt{n}$ — — Weak $1 < \alpha \leq 1 + \sqrt{y}$ — $(1 + \sqrt{y})^2$ — 0 — Edge / Tracy-Widom, $p^{-2/3}$ —

(b) One-paragraph synthesis: the BBP phase transition is the cleanest quantitative example of when high-dimensional inference recovers structure and when it does not. It explains, in a controlled toy model, why methods like PCA and spectral clustering succeed only when signal exceeds a threshold determined by the aspect ratio.

(c) Outlook: open questions to flag. - The critical case $\alpha = 1 + \sqrt{y}$. The current literature does not yet contain a clean limit theorem; cite open work. - Exact variance $\sigma_{\alpha, y}^2$ in the strong-spike CLT. - Beyond Johnstone's case: general spiked models with non-flat bulk (cite Baik and Silverstein (2006); Benaych-Georges and Nadakuditi (2011)). - Rigorous spectral clustering derivations beyond the toy model; cite Couillet and Benaych-Georges (2016).

(d) Closing sentence: the goal of this thesis was to develop the foundational picture from one end of the dimension spectrum to the other; the open questions above point to where the field is going.

A Auxiliary Results

[TODO] Content to fill in:

A.1 Borel-Cantelli Lemmas. - First Borel-Cantelli: if $\sum_n \mathbb{P}(A_n) < \infty$, then $\mathbb{P}(A_n \text{ i.o.}) = 0$. - Second Borel-Cantelli (independence version): if events are independent and $\sum_n \mathbb{P}(A_n) = \infty$, then $\mathbb{P}(A_n \text{ i.o.}) = 1$. - Used in Section 2.4 for the convergence-in-probability but not-almost-surely counterexample.

A.2 Slutsky's theorem. - Statement and a one-paragraph proof or reference to Casella and Berger (2002, Thm. 5.5.17).

A.3 Continuous mapping theorem. - Statement: if $X_n \xrightarrow{d} X$ and g is continuous on a set of $\mathbb{P}_X$ -probability one, then $g(X_n) \xrightarrow{d} g(X)$.

A.4 Multivariate delta method. - Statement and reference; used in Sections 3.2 and 4.3.

A.5 Characteristic functions and the Lévy continuity theorem. - Definition: $\varphi_X(t) = \mathbb{E}[e^{itX}]$; always exists. - Examples: $\mathcal{N}(0, 1)$ gives $e^{-t^2/2}$, Cauchy gives $e^{-|t|}$. - Lévy: $X_n \xrightarrow{d} X \iff \varphi_{X_n}(t) \rightarrow \varphi_X(t)$ pointwise, with continuity at 0 sufficient for the converse. - Proof of the classical CLT via $(1 + t/n)^n \rightarrow e^t$ expansion.

B Simulation Methodology

[TODO] Content to fill in:

B.1 General reproducibility conventions. - All notebooks use Python 3, with `numpy`, `matplotlib`, and `scipy.stats`. - Random number generation: `rng = np.random.default_rng(seed)`, with explicit `seed` at the top of every notebook. - Eigendecomposition: `np.linalg.eigh` on symmetric inputs; the top eigenpair is taken as `evals[-1]`, `evecs[:, -1]`. - Plot styling: shared helper module `notebooks/thesis_style.py` applies the six-color palette and sans-serif math labels.

B.2 Per-notebook parameter table. - Notebook 01 (LLN/CLT demos): $n \in \{100, 1000, 10000\}$, 500 trials. - Notebook 02 (Wishart and repulsion): $n = 100$, $p = 2$, 5000 trials. - Notebook 03 (fixed-dim fluctuations): $n = 500$, 5000 trials, both $\Sigma = I_2$ and $\Sigma = \text{diag}(1, 2)$. - Notebook 04 (Marchenko-Pastur & TW): $n = 1000$, $y = 1/2$, 1 trial for bulk density, 1000 trials for top eigenvalue. - Notebook 05 (spiked eigenvalues, BBP): $y = 1/2$, $n \in \{100, 200, 400, 800, 1200\}$, 200 trials per n , $\alpha \in \{1.4, 2.5\}$. - Notebook 06 (fluctuation Q-Q): $n = 1200$, $y = 1/2$, 500 trials, $\alpha = 2.5$. - Notebook 07 (eigenvector overlap): $n = 1200$, $y = 1/2$, 500 trials, $\alpha \in \{1.4, 2.5\}$. - Notebook 08 (spectral clustering toy): two-class Gaussian mixture, $n = 600$ per class, $p \in \{50, 200\}$.

B.3 Numerical caveats. - For the strong-spike fluctuation Q-Q plot, we standardize empirically (subtract sample mean, divide by sample standard deviation) because the exact variance $\sigma_{\alpha, y}^2$ requires constants from Yao et al. (2015, Thm. 11.11). - For the weak spike, we do not attempt an exact Tracy-Widom fit because the finite-sample centering and scaling constants are delicate; we report the qualitative shape of the histogram and defer the rigorous fit to future work. - The realized aspect ratio p/n differs from the target y when yn is not integer; figures use the realized ratio in the bulk-edge formula $(1 + \sqrt{p/n})^2$ for consistency.

C Selected Code Listings

[TODO] *Include 4 to 6 of the most informative snippets from the notebooks (not full notebooks). Each listing should be self-contained at the function level, captioned, and referenced from the main text where the corresponding figure appears.*

Suggested listings to include here:

- *C.1 One-spike Johnstone simulation (driving function for Section 4.2). Source: `notebooks/05_spiked`*
- *C.2 Eigenvector overlap simulation (Section 4.4). Source: `notebooks/07_eigenvector_overlap.ipynb`*
- *C.3 Strong-spike fluctuation simulation (Section 4.3). Source: `notebooks/06_fluctuation_qq_plots`*
- *C.4 Marchenko-Pastur bulk simulation (Section 3.3). Source: `notebooks/04_marchenko_pastur_trace`*

Each listing will be wrapped in a `lstlisting` environment using the `thesisPython` style defined in `style/thesis.sty`, with a caption and a `label` for cross-referencing from the main text.

References

- Jinho Baik and Jack W. Silverstein. Eigenvalues of large sample covariance matrices of spiked population models. *Journal of Multivariate Analysis*, 97(6):1382–1408, 2006.
- Jinho Baik, Gérard Ben Arous, and Sandrine Péché. Phase transition of the largest eigenvalue for nonnull complex sample covariance matrices. *The Annals of Probability*, 33(5):1643–1697, 2005.
- Florent Benaych-Georges and Raj Rao Nadakuditi. The eigenvalues and eigenvectors of finite, low rank perturbations of large random matrices. *Advances in Mathematics*, 227(1):494–521, 2011.
- Patrick Billingsley. *Probability and Measure*. Wiley, 3rd edition, 1995.
- George Casella and Roger L. Berger. *Statistical Inference*. Duxbury, 2nd edition, 2002.
- Romain Couillet and Florent Benaych-Georges. Kernel spectral clustering of large dimensional data. *Electronic Journal of Statistics*, 10(1):1393–1454, 2016.
- Romain Couillet and Zhenyu Liao. Random matrix methods for machine learning, 2022. Cambridge University Press. Online draft: <https://zhenyu-liao.github.io/pdf/RMT4ML.pdf>.
- Delphine Féral and Sandrine Péché. The largest eigenvalues of sample covariance matrices for a spiked population: diagonal case. *Journal of Mathematical Physics*, 50(7):073302, 2009.
- Joseph Genzer. Spectral analysis of multiscale sample covariance matrices, 2025. Honors Thesis, Tulane University.
- Iain M. Johnstone. On the distribution of the largest eigenvalue in principal components analysis. *The Annals of Statistics*, 29(2):295–327, 2001.
- V. A. Marchenko and L. A. Pastur. Distribution of eigenvalues for some sets of random matrices. *Mathematics of the USSR-Sbornik*, 1(4):457–483, 1967.
- Andrew Y. Ng, Michael I. Jordan, and Yair Weiss. On spectral clustering: Analysis and an algorithm. In *Advances in Neural Information Processing Systems (NeurIPS) 14*, pages 849–856, 2002.
- Craig A. Tracy and Harold Widom. On orthogonal and symplectic matrix ensembles. *Communications in Mathematical Physics*, 177(3):727–754, 1996.
- Roman Vershynin. *High-Dimensional Probability: An Introduction with Applications in Data Science*. Cambridge University Press, 2018.
- Ulrike von Luxburg. A tutorial on spectral clustering. *Statistics and Computing*, 17(4):395–416, 2007.
- Jianfeng Yao, Shurong Zheng, and Zhidong Bai. *Large Sample Covariance Matrices and High-Dimensional Data Analysis*. Cambridge University Press, 2015.